

Concurrency: Multi-core Programming & Data Processing



Mutual Exclusion

Prof. P. Felber

Pascal.Felber@unine.ch

<http://iiun.unine.ch/>

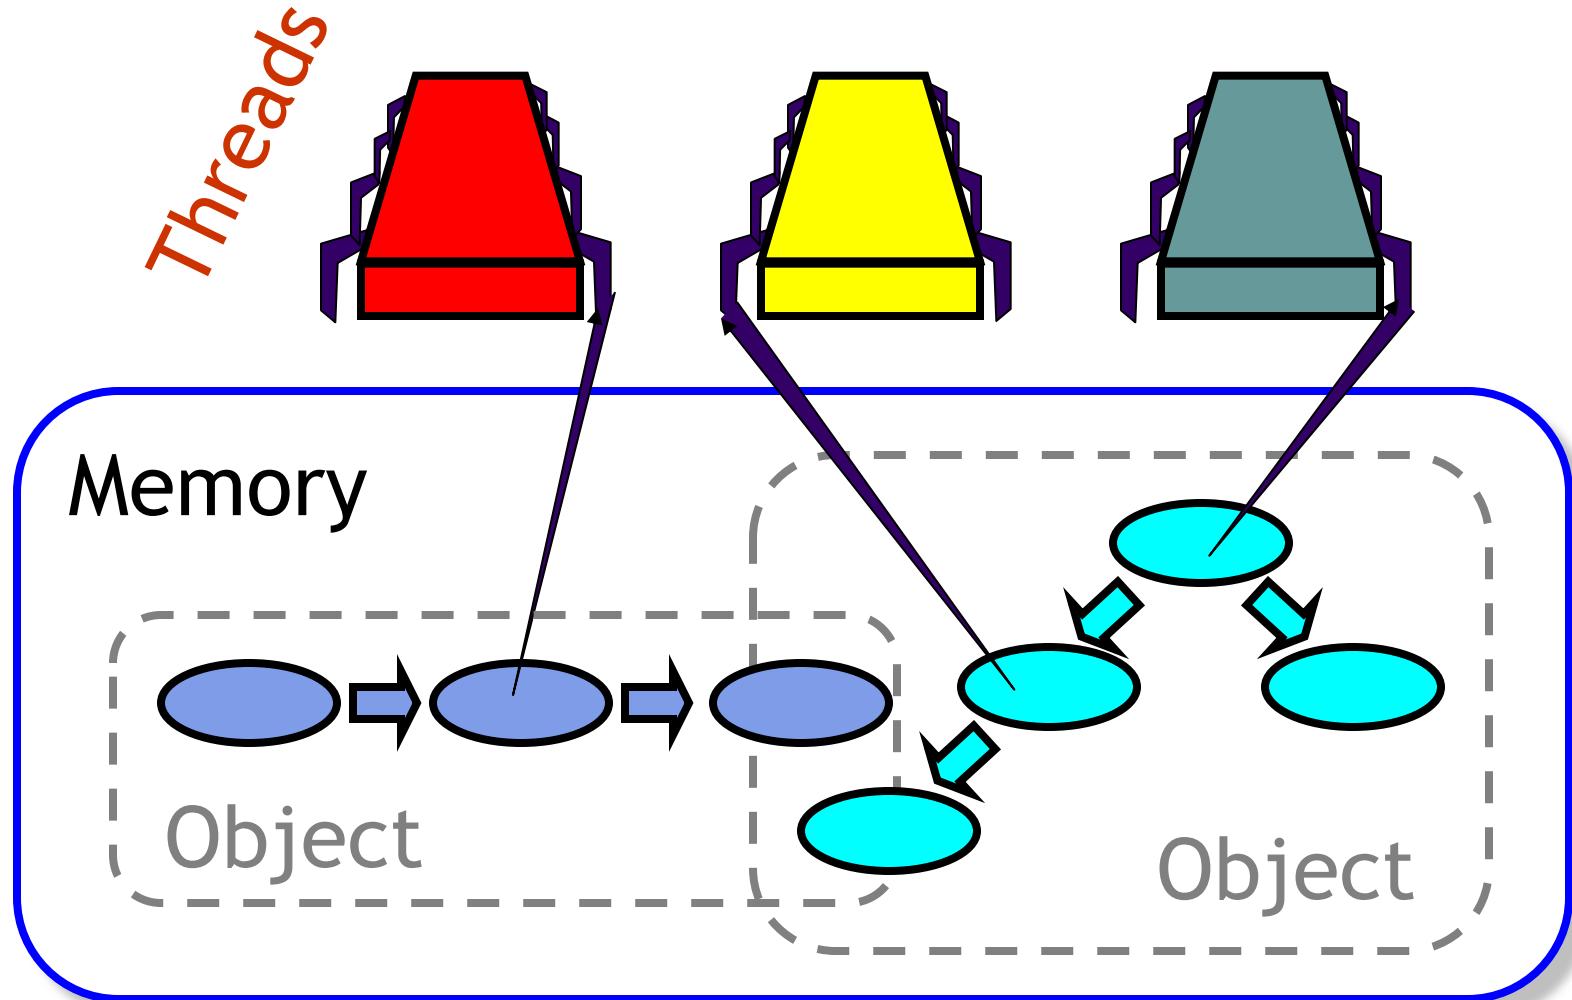
Based on slides by Maurice Herlihy and Nir Shavit



Review

- Model of computation
- Asynchrony
- Mutual exclusion

Concurrent Computation



Asynchrony

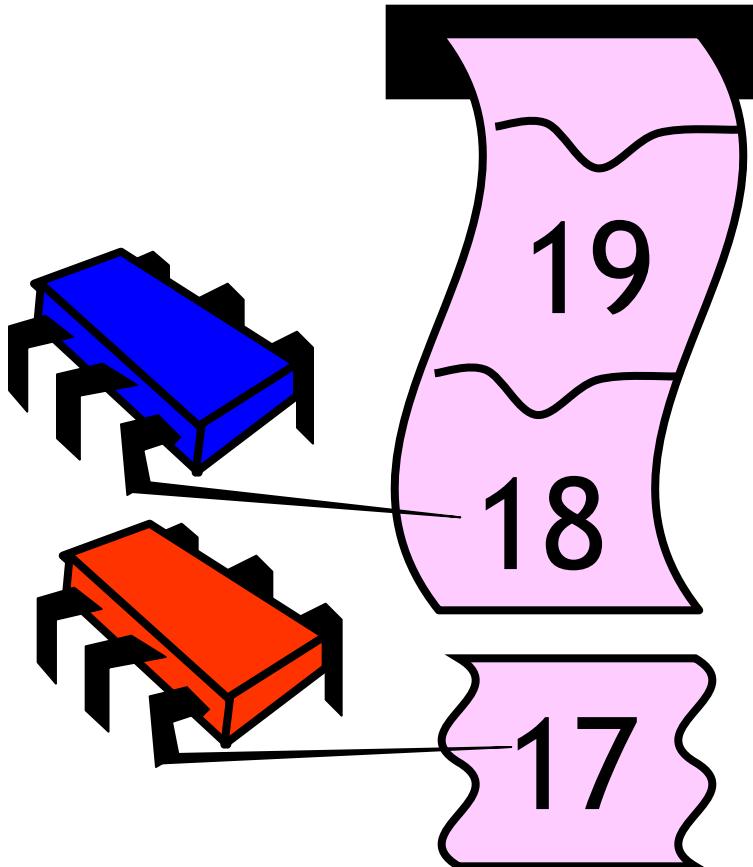


- Sudden unpredictable delays
 - Cache misses (*short*)
 - Page faults (*long*)
 - Scheduling quantum used up (*really long*)

Parallel Primality Testing

- Challenge
 - Print primes from 1 to 10^{10}
- Given
 - Ten-processor multiprocessor
 - One thread per processor
- Goal
 - Get ten-fold speedup (or close)

Shared Counter



Each thread
takes a number

Procedure for Thread i

```
Counter counter = new Counter(1);

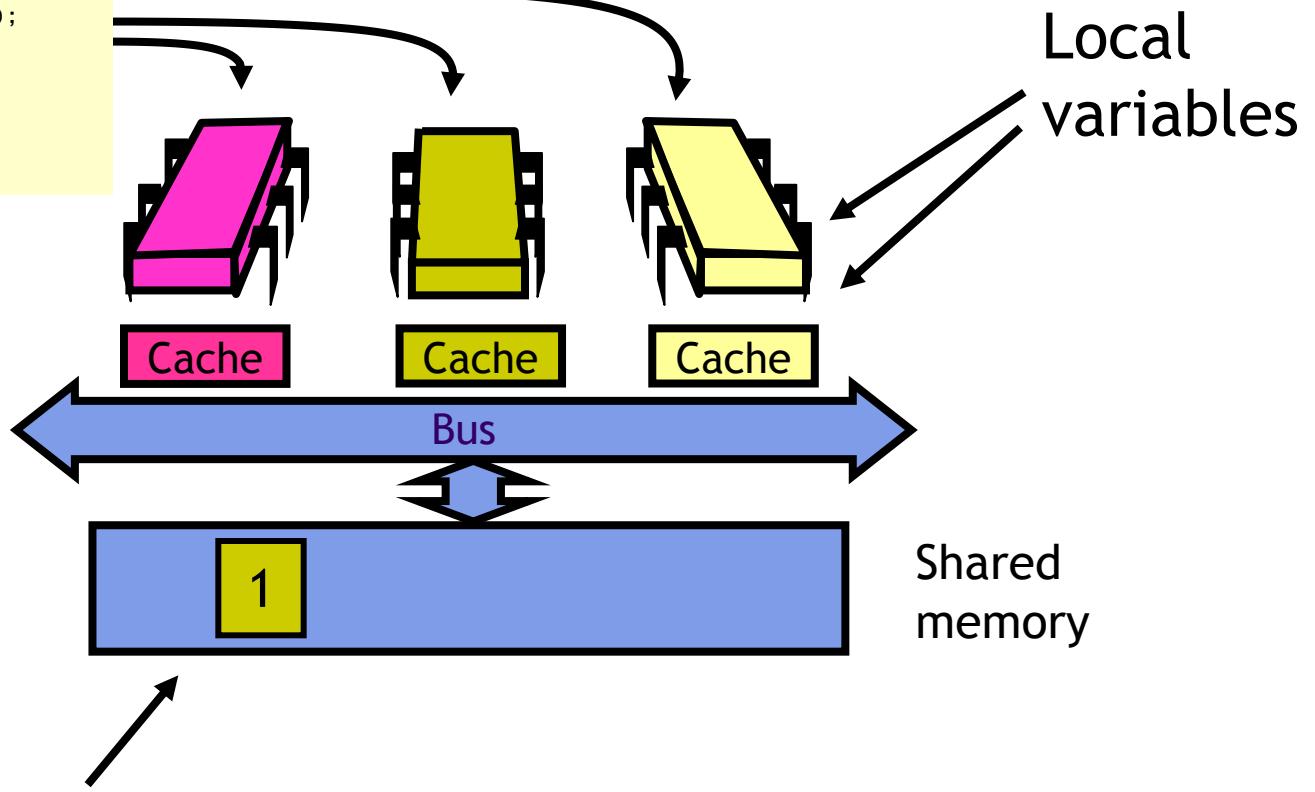
void primePrint {
    while (true) {
        long j = counter.getAndIncrement();
        if (j > 1010)
            break;
        if (isPrime(j))
            print(j);
    }
}
```

Where Things Reside

```
Counter counter = new Counter(1);
void primePrint {
    while (true) {
        long j = counter.getAndIncrement();
        if (j > 1010)
            break;
        if (isPrime(j))
            print(j);
    }
}
```

Code

Shared counter



Challenge

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        long temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```

Make these steps
atomic (indivisible)

Mutual Exclusion



- Today we will try to formalize our understanding of mutual exclusion
- We will also use the opportunity to show you how to argue about and prove various properties in an asynchronous concurrent setting

Mutual Exclusion

- Formal problem definitions
- Solutions for 2 threads
- Solutions for n threads
- Fair solutions
- Inherent costs



Warning

- You will never use these protocols
 - Get over it
- You are advised to understand them
 - The same issues show up everywhere
 - Except hidden and more complex

Why is Concurrent Programming so Hard?

- Try preparing a seven-course banquet
 - By yourself
 - With one friend
 - With twenty-seven friends...
- Before we can talk about programs
 - Need a language
 - Describing time and concurrency

What is Time?

- “*Absolute, true and mathematical time, of itself and from its own nature, flows equably without relation to anything external.*” (I. Newton, 1689)
- “*Time is, like, Nature’s way of making sure that everything doesn’t happen all at once.*” (Anonymous, circa 1968)

Time



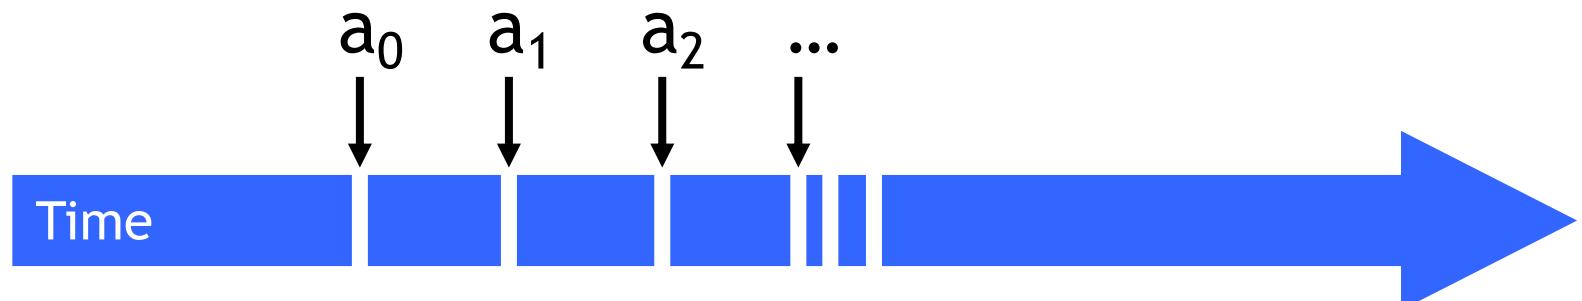
Events

- An **event** a_0 of thread A is
 - Instantaneous
 - No simultaneous events (break ties)



Threads

- A **thread** A is (formally) a sequence a_0, a_1, \dots of events
 - “Trace” model
 - Notation: $a_0 \rightarrow a_1$ indicates order

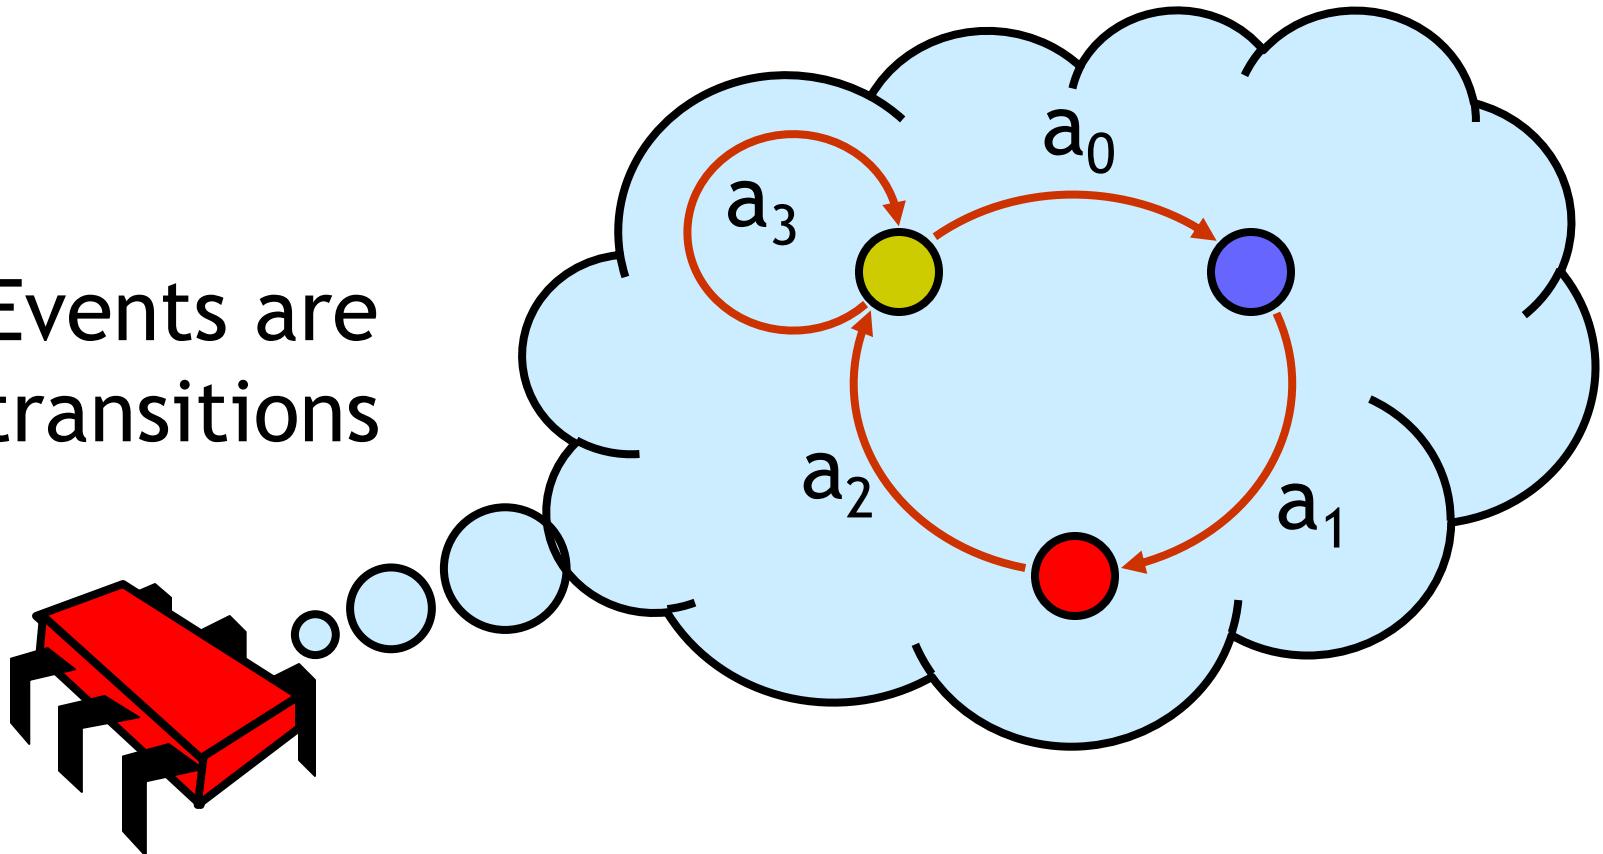


Example Thread Events

- Assign to shared variable
- Assign to local variable
- Invoke method
- Return from method
- Lots of other things...

Threads are State Machines

Events are transitions



States

- Thread state
 - Program counter
 - Local variables
- System state
 - Object fields (shared variables)
 - Union of thread states

Concurrency

- Thread A

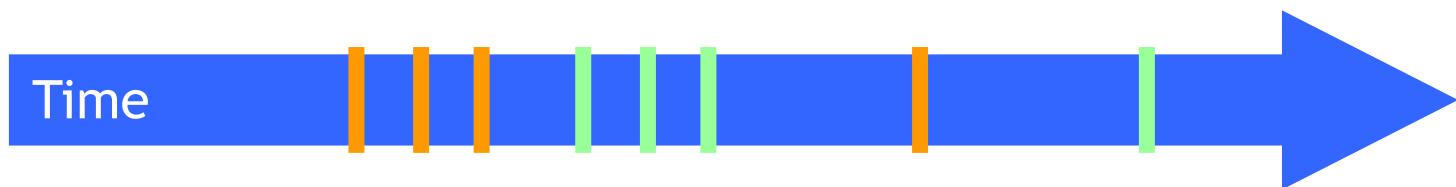


- Thread B



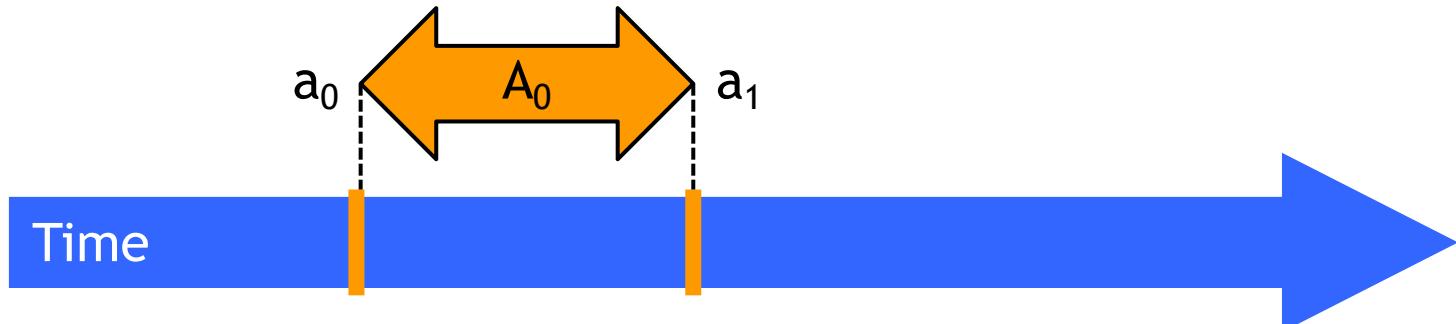
Interleavings

- Events of two or more threads
 - Interleaved
 - Not necessarily independent (why?)

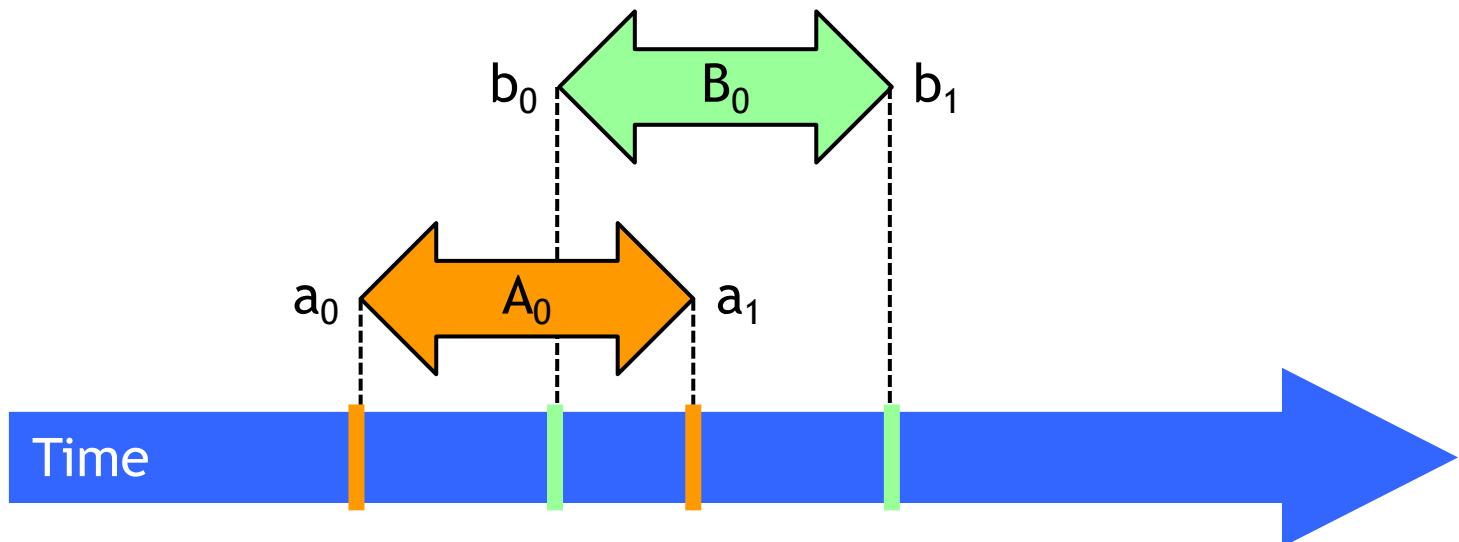


Intervals

- An **interval** $A_0 = (a_0, a_1)$ is
 - Time between events a_0 and a_1

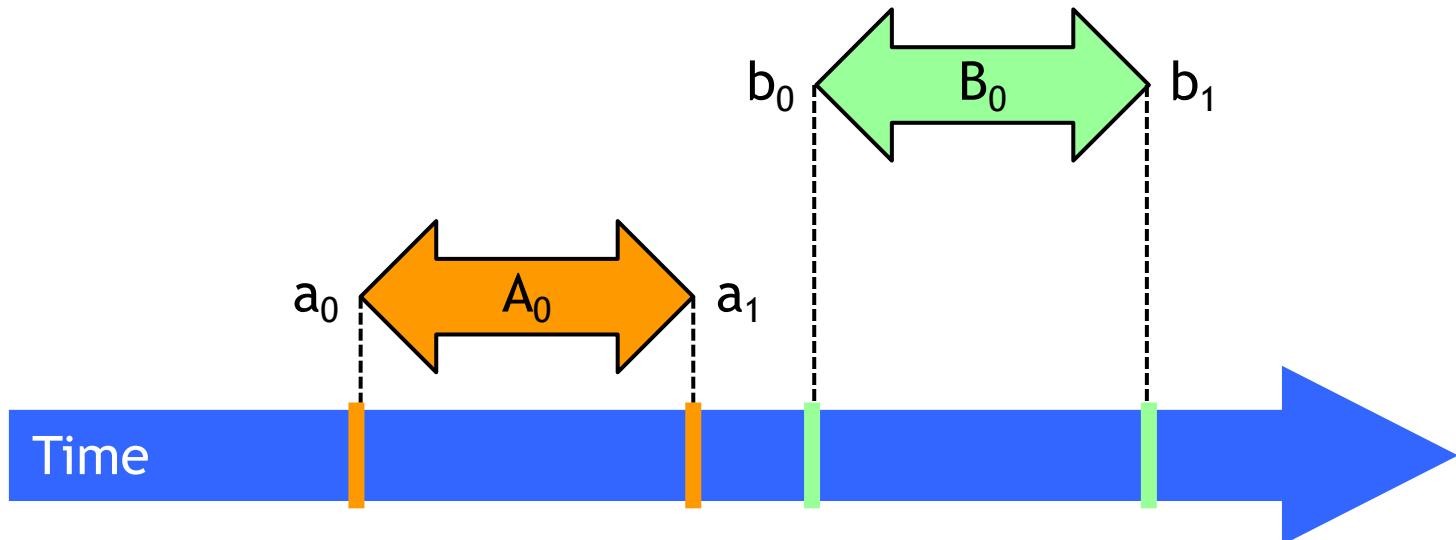


Intervals may Overlap



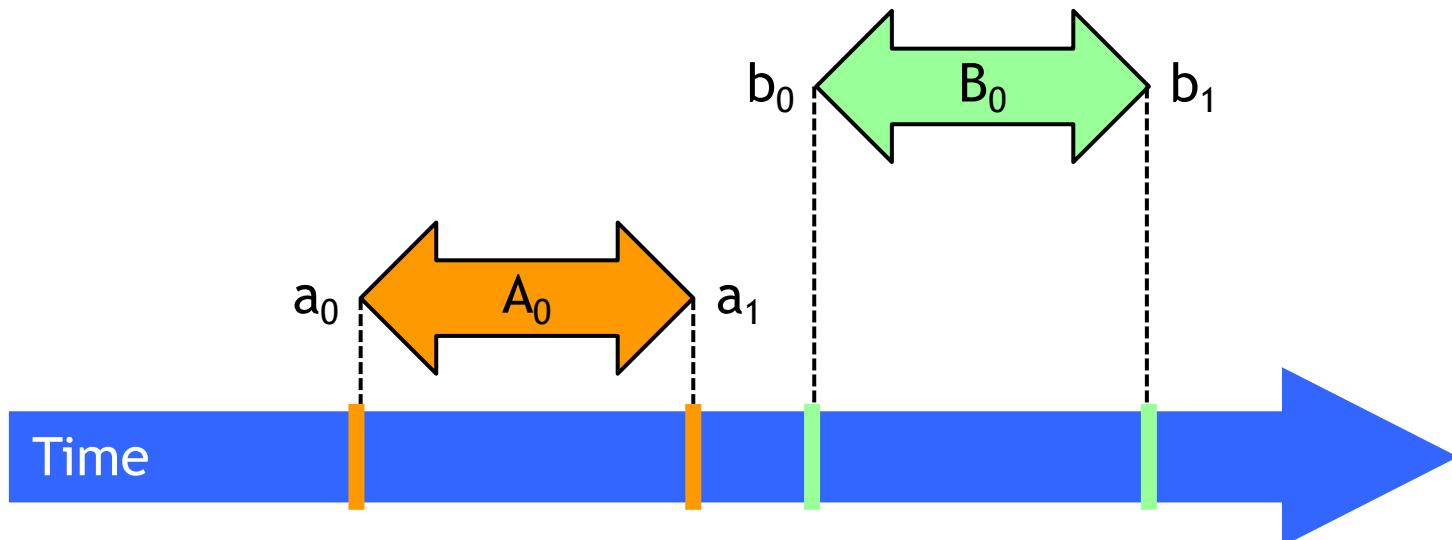
Intervals may be Disjoint

- Precedence
 - Interval A_0 **precedes** interval B_0



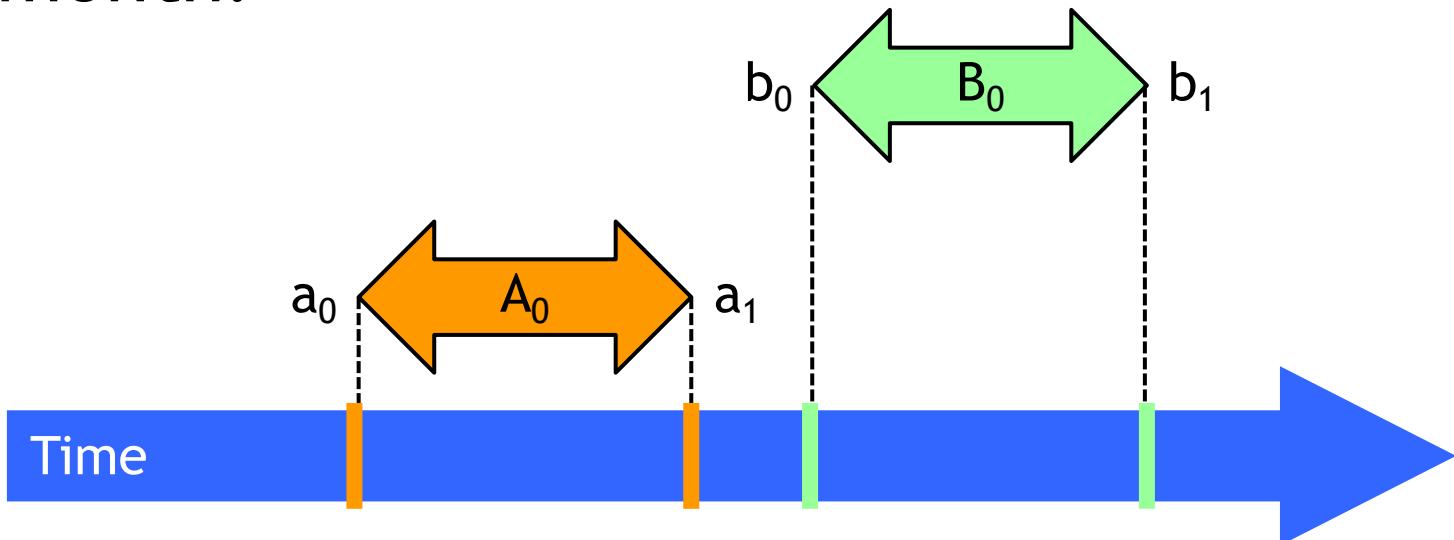
Precedence

- Notation: $A_0 \rightarrow B_0$
- Formally:
 - End event of A_0 before start event of B_0
 - Also called “happens before” or “precedes”



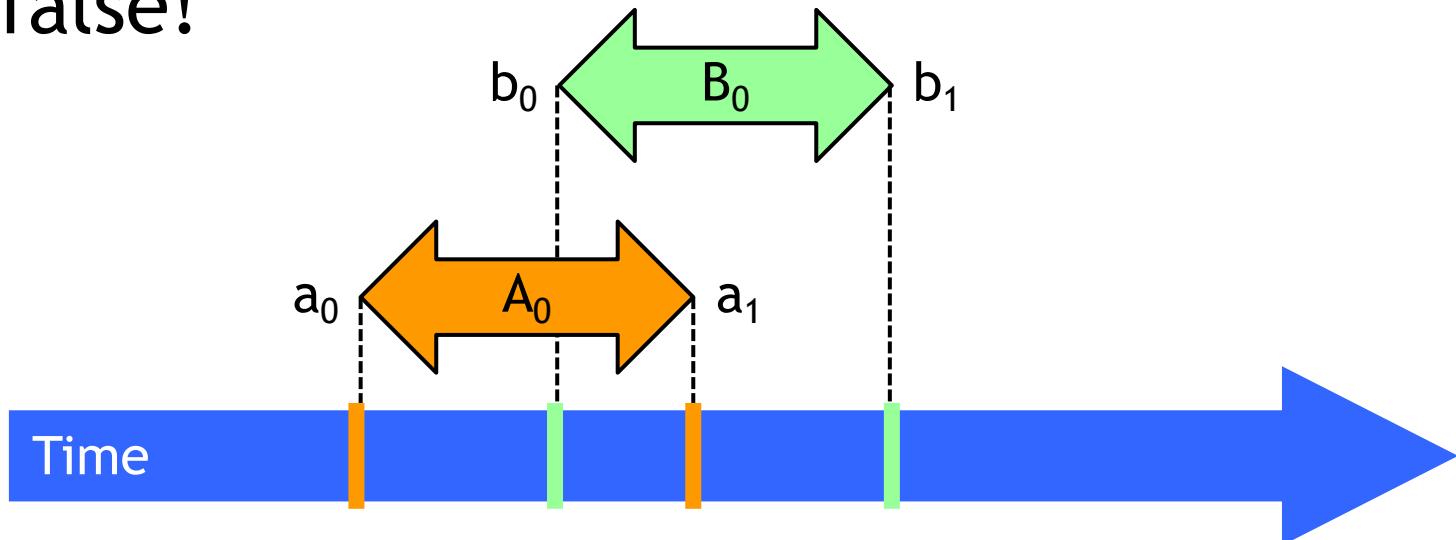
Precedence Ordering

- Remark: $A_0 \rightarrow B_0$ is just like saying
 - 1066 AD \rightarrow 1492 AD
 - Middle Ages \rightarrow Renaissance
- Oh wait, what about this week vs. this month?



Precedence Ordering

- Never true that $A \rightarrow A$
- If $A \rightarrow B$ then not true that $B \rightarrow A$
- If $A \rightarrow B$ and $B \rightarrow C$ then $A \rightarrow C$
- Funny thing: $A \rightarrow B$ and $B \rightarrow A$ might both be false!



Partial Orders

(you may know this already)

- **Irreflexive**

- Never true that $A \rightarrow A$

- **Antisymmetric**

- If $A \rightarrow B$ then not true that $B \rightarrow A$

- **Transitive**

- If $A \rightarrow B$ and $B \rightarrow C$ then $A \rightarrow C$

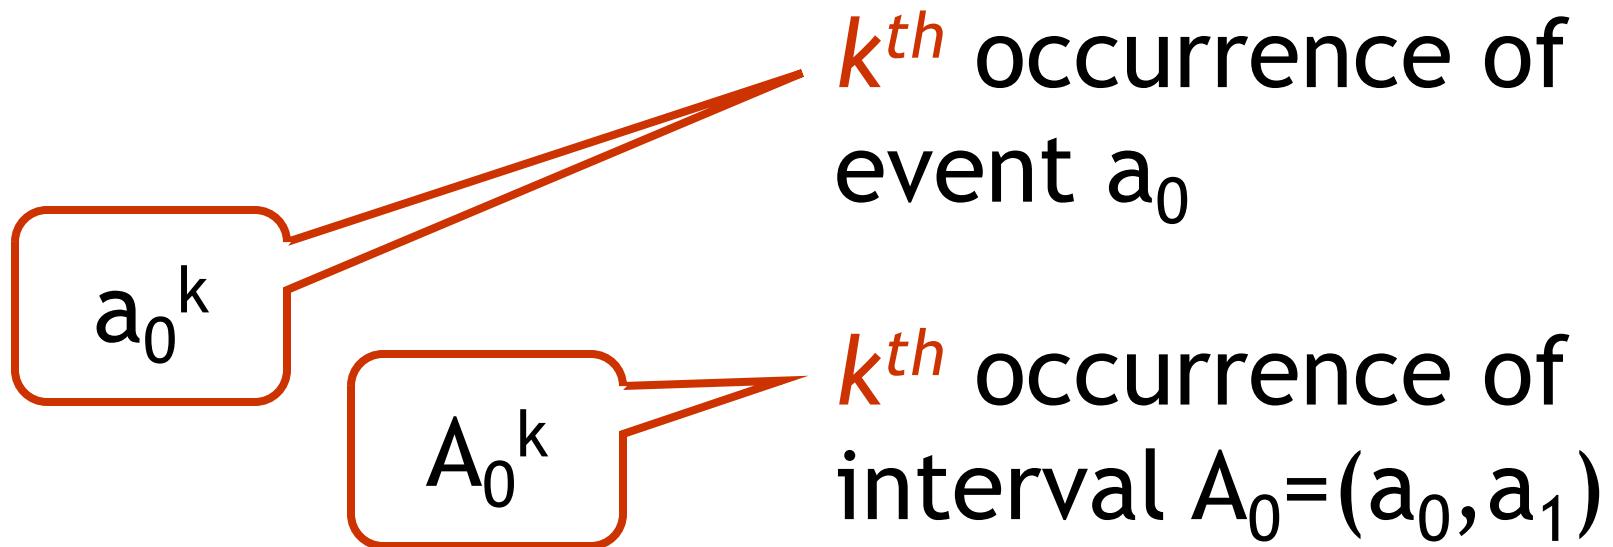
Total Orders

(you may know this already)

- Also
 - Irreflexive
 - Antisymmetric
 - Transitive
- Except that for every distinct A, B
 - Either $A \rightarrow B$ or $B \rightarrow A$

Repeated Events

```
while (mumble) {  
    a0; a1;  
}
```



Implementing a Counter

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        long temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```

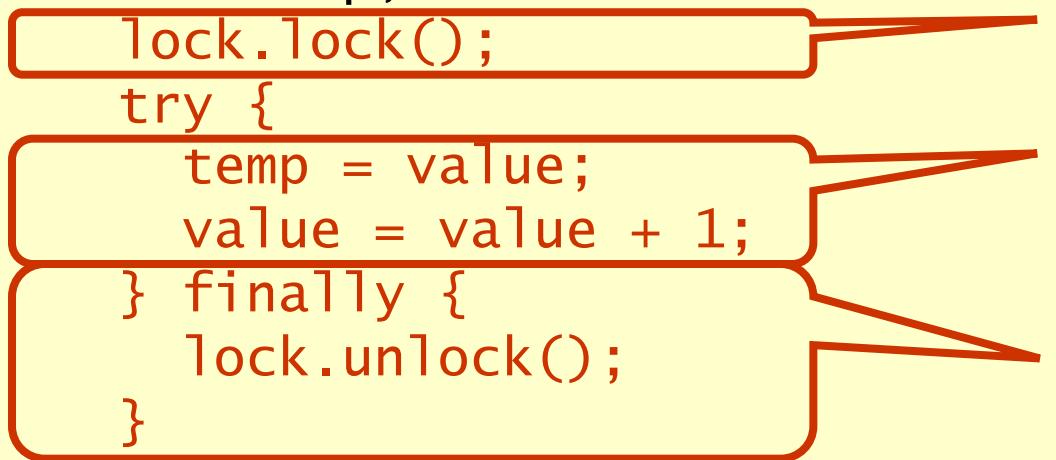
Make these steps
atomic (indivisible)

Locks (Mutual Exclusion)

```
public interface Lock {  
    public void lock(); // Acquire lock  
    public void unlock(); // Release lock  
}
```

Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        int temp;  
        lock.lock(); // Acquire lock  
        try {  
            temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock(); // Release lock  
        }  
        return temp;  
    }  
}
```



The diagram illustrates the use of locks in the `getAndIncrement()` method. It uses red boxes and arrows to highlight specific sections of the code:

- A red box encloses the line `lock.lock();`, with an arrow pointing to the text "Acquire lock".
- A red box encloses the block of code between `try {` and `} finally {` (excluding the finally block itself), with an arrow pointing to the text "Critical section".
- A red box encloses the line `lock.unlock();` in the finally block, with an arrow pointing to the text "Release lock (no matter what)".

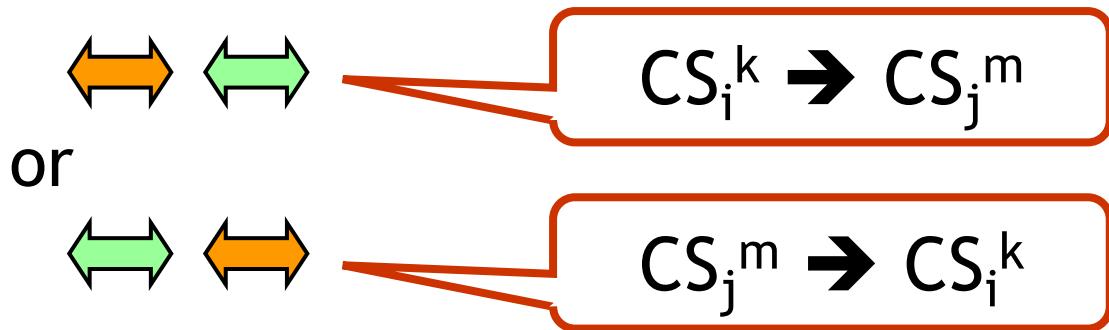
Using Locks

```
private Lock lock;  
...  
while{mumble}{  
    <non critical section>  
    lock.lock();  
    <critical section>  
    lock.unlock();  
}  
...
```

- Only **well formed executions!**
- What are the required properties of locks?

Mutual Exclusion

- Let $CS_i^k \leftrightarrow$ be thread i's k^{th} critical section execution
- And $CS_j^m \leftrightarrow$ thread j's m^{th} critical section execution
- Then either



Deadlock-Free



- If some thread calls **lock()**
 - And never returns
 - Then other threads must complete **lock()** and **unlock()** calls infinitely often
- System as a whole makes progress
 - Even if individuals starve

Lockout-Free

- If some thread calls **lock()**
 - It will eventually return
- Individual threads make progress



Two-Thread vs. n -Thread

- Two-thread solutions first
 - Illustrate most basic ideas
 - Fits on one slide
- Then n -thread solutions

Two-Thread Conventions

```
class ... implements Lock {  
    ...  
    // Thread-local index, 0 or 1  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
        ...  
    }  
}
```

Henceforth: **i** is current thread, **j** is other thread

LockOne

```
class Lockone implements Lock {  
    private volatile boolean[]  
        flag = new boolean[2];  
  
    public void lock() {  
        flag[i] = true;           Set my flag  
        while (flag[j]) {}      Wait for other  
    }                          flag to go false  
  
    public void unlock() {  
        flag[i] = false;         No longer  
    }                          interested  
}
```

LockOne Satisfies Mutual Exclusion

- Assume CS_A^j overlaps CS_B^k
- Consider each thread's (j^{th} and k^{th}) last read and write in the **lock()** method before entering
- Derive a contradiction

From the Code

$\text{write}_A(\text{flag}[A]=\text{true}) \rightarrow \text{read}_A(\text{flag}[B]==\text{false}) \rightarrow CS_A$

$\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow \text{read}_B(\text{flag}[A]==\text{false}) \rightarrow CS_B$

```
class Lockone implements Lock {  
    private volatile boolean[] flag = new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

From the Assumption

$\text{read}_A(\text{flag}[B]==\text{false}) \rightarrow \text{write}_B(\text{flag}[B]=\text{true})$

$\text{read}_B(\text{flag}[A]==\text{false}) \rightarrow \text{write}_A(\text{flag}[A]=\text{true})$

Combining

Assumptions

$\text{read}_A(\text{flag}[B]==\text{false}) \rightarrow \text{write}_B(\text{flag}[B]=\text{true})$

$\text{read}_B(\text{flag}[A]==\text{false}) \rightarrow \text{write}_A(\text{flag}[A]=\text{true})$

From the code

$\text{write}_A(\text{flag}[A]=\text{true}) \rightarrow \text{read}_A(\text{flag}[B]==\text{false})$

$\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow \text{read}_B(\text{flag}[A]==\text{false})$

Cycle!

Deadlock Freedom

- LockOne fails deadlock-freedom
 - Concurrent execution can deadlock

Time ↓

```
flag[i] = true;  
while (flag[j]) {}
```

```
flag[j] = true;  
while (flag[i]) {}
```

- Sequential executions OK

LockTwo

```
public class LockTwo implements Lock {  
    private volatile int victim;  
    public void lock() {  
        victim = i; Let other  
go first  
        while (victim == i) {} Wait for  
permission  
    }  
    public void unlock() {} Nothing  
to do  
}
```

LockTwo Claims

- Satisfies mutual exclusion

- If thread **i** in CS
- Then **victim == j**
- Cannot be both **i** and **j**

```
{  
    victim = i;  
    while (victim == i) {}  
}
```

- Not deadlock free
 - Sequential execution deadlocks
 - Concurrent execution does not

Peterson's Algorithm

```

public void lock() {
    flag[i] = true;           Announce I'm interested
    victim = i;              Defer to other
    while (flag[j] && victim == i) {}
}
}   Wait while other interested & I'm the victim
public void unlock() {
    flag[i] = false;         No longer interested
}

```

Mutual Exclusion

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}
```

- If thread **0** in critical section
 - **flag[0] == true**
 - **victim == 1**
- If thread **1** in critical section
 - **flag[1] == true**
 - **victim == 0**

Cannot both be true

Deadlock Free

```
public void lock() {  
    ...  
    while (flag[j] && victim == i) {}  
}
```

- Thread blocked
 - Only at **while** loop
 - Only if it is the victim
- One or the other must not be the victim

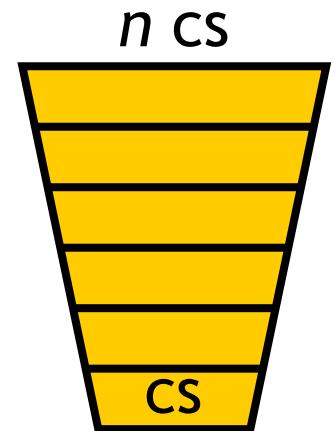
Lockout Free

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}  
public void unlock() {  
    flag[i] = false;  
}
```

- Thread **i** blocked only if **j** repeatedly re-enters with **flag[j] == true** and **victim == i**
- When **j** re-enters, it sets **victim** to **j**, so **i** gets in

The Filter Algorithm for n Threads

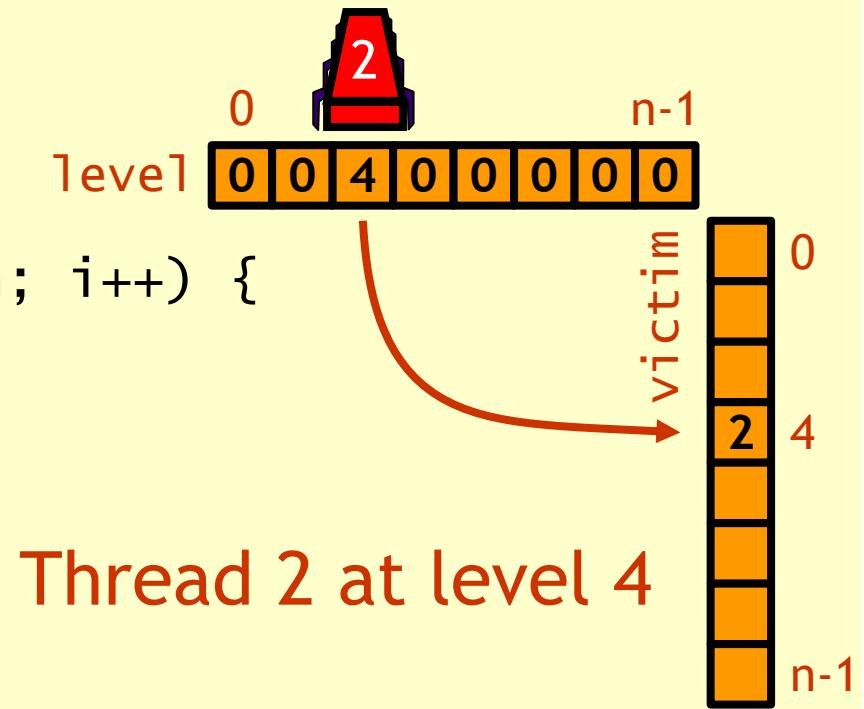
- There are $n-1$ “waiting rooms” called levels
- At each level
 - At least one enters level
 - At least one blocked if many try
- Only one thread makes it through



Filter

```
class Filter implements Lock {
    volatile int[] level; // level[i] for thread i
    volatile int[] victim; // victim[L] for level L

    public Filter(int n) {
        level = new int[n];
        victim = new int[n];
        for (int i = 0; i < n; i++) {
            level[i] = 0;
        }
    }
    ...
}
```



Filter

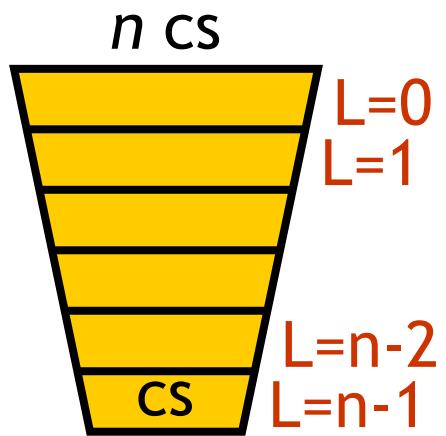
```

class Filter implements Lock {
  ...
  public void lock() {
    for (int L = 1; L < n; L++) { One level at a time
      level[i] = L; Announce intention
      victim[L] = i; to enter level L
    }
    while (( $\exists$ k != i | level[k] >= L) &&
           victim[L] == i) {} Give priority to anyone but me
  } Wait as
  } long as
  someone else is at same or higher level and I'm
  designated victim (enter level L when loop
  public void unlock() { level[i] = 0; } complete)
}

```

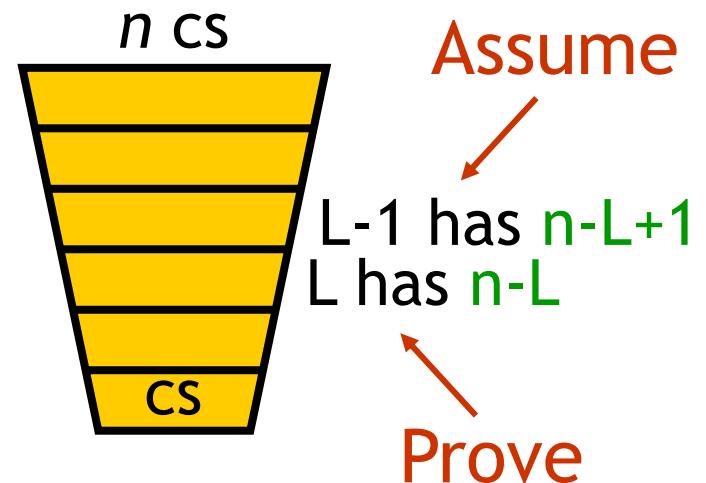
Claim

- Start at level $L=0$
- At most $n-L$ threads enter level L
- Mutual exclusion at level $L=n-1$

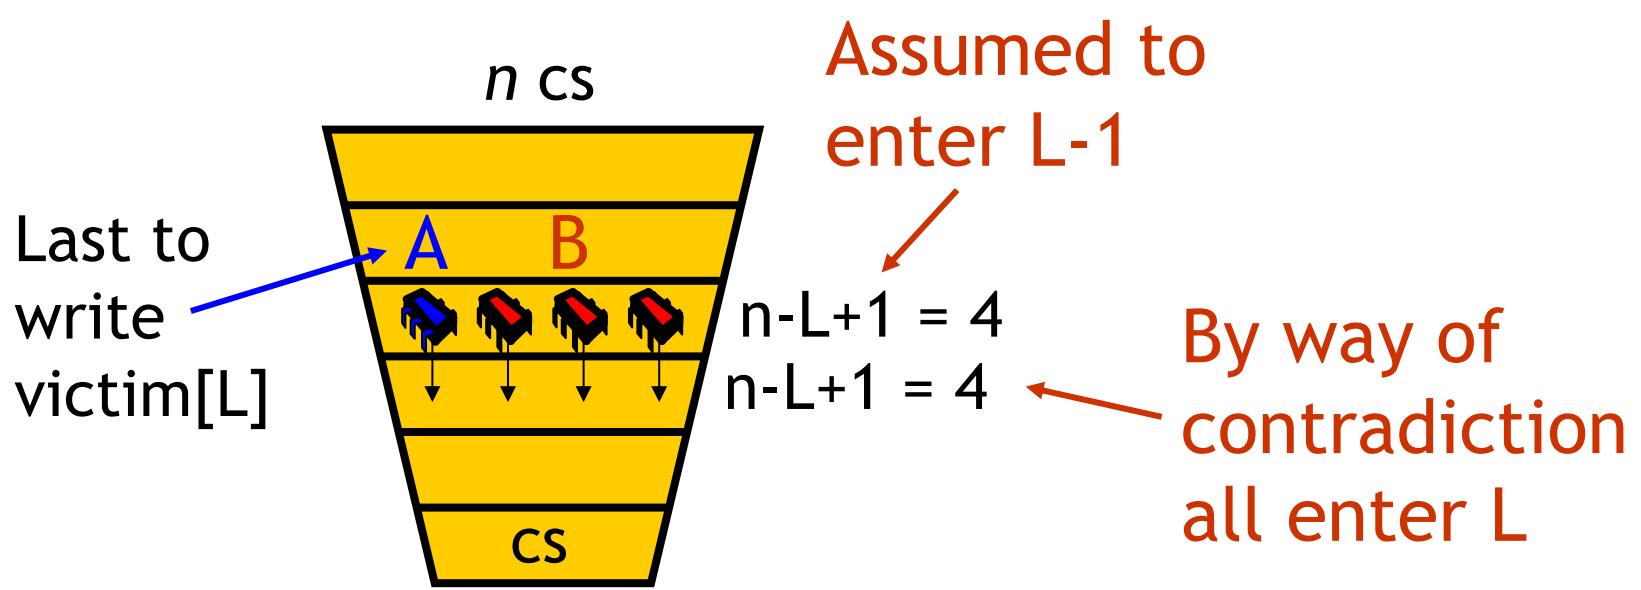


Induction Hypothesis

- No more than $n-L+1$ threads at level $L-1$
- Induction step: by contradiction
- Assume all at level $L-1$ enter level L
- A last to write **victim[L]**
- B is any other thread at level L



Proof Structure



- Show that A must have seen B at level L and since $\text{victim}[L] == A$ could not have entered

Observations

1. $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$ (code)
2. $\text{write}_A(\text{victim}[L]=A) \rightarrow \text{read}_A(\text{level}[B])$ (code)
3. $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$
 (by assumption: A is the last thread to write victim[L])

```

public void lock() {
    for (int L = 1; L < n; L++) {
        level[i] = L;
        victim[L] = i;
        while ((∃k != i | level[k] >= L) && victim[L] == i) {}
    }
}
  
```

Combining Observations

1. $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$ (code)
3. $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$ (ass.)
2. $\text{write}_A(\text{victim}[L]=A) \rightarrow \text{read}_A(\text{level}[B])$ (code)

Thus, A read $\text{level}[B] \geq L$, A was last to write $\text{victim}[L]$, so it could not have entered level L!

```
public void lock() {
    for (int L = 1; L < n; L++) {
        level[i] = L;
        victim[L] = i;
        while ((∃k != i | level[k] >= L) && victim[L] == i) {}
    }
}
```

No Lockout

- Filter Lock satisfies “no lockout” properties
 - Just like Peterson algorithm at any level
 - So no one starves (no lockout)
- But what about fairness?
 - Threads can be overtaken by others

Bounded Waiting

- Want stronger fairness guarantees
- Thread not “overtaken” too much
- Need to adjust definitions...
- Divide **Lock()** method into 2 parts
 - Doorway interval (D_A), always finishes in finite number of steps
 - Waiting interval (W_A), may take unbounded number of steps

r-Bounded Waiting

- For threads A and B
 - If $D_A^k \rightarrow D_B^j$
 - A's k^{th} doorway precedes B's j^{th} doorway
 - Then $CS_A^k \rightarrow CS_B^{j+r}$
 - A's k^{th} critical section precedes B's $(j+r)^{\text{th}}$ critical section
 - B cannot overtake A by more than r times
- First-come-first-served if $r = 0$

Fairness Again

- Filter Lock satisfies some properties
 - No one starves (no lockout)
 - But very weak fairness
 - Not r -bounded for any r !
 - That's pretty lame...

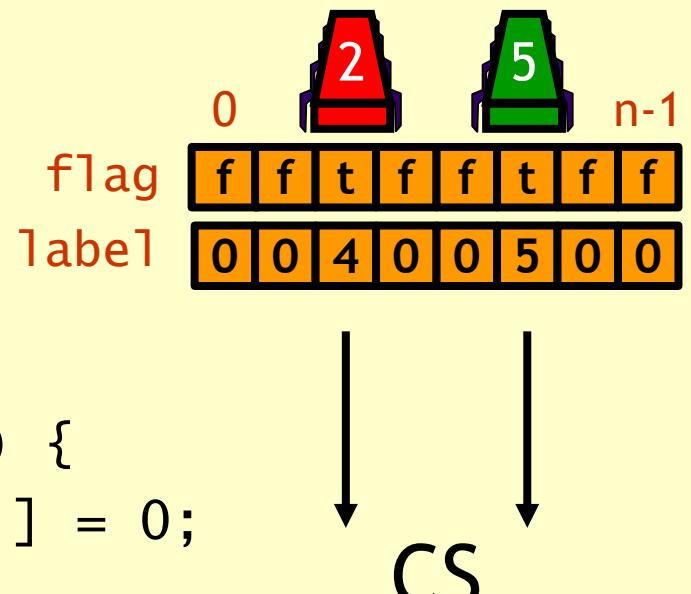
Bakery Algorithm

- Provides First-Come-First-Served
- How?
 - Take a “number”
 - Wait until lower numbers have been served
- Lexicographic order
 - $(a,i) > (b,j)$
 - If $a > b$, or $a == b$ and $i > j$

Bakery Algorithm

```

class Bakery implements Lock {
    volatile boolean[] flag;
    volatile int[] label;
    public Bakery (int n) {
        flag = new boolean[n];
        label = new int[n];
        for (int i = 0; i < n; i++) {
            flag[i] = false; label[i] = 0;
        }
    }
    ...
}
  
```



Bakery Algorithm

Doorway

```

class Bakery implements Lock {
  ...
  public void lock() {
    flag[i] = true;           Take increasing label
    label[i] = max(label[0],...,label[n-1])+1;   (read labels in some
                                                arbitrary order)
    while (( $\exists k \mid \text{flag}[k]$ ) && ( $\text{label}[i], i > \text{label}[k], k$ )) {}
  }
  ...
}
  
```

I'm interested

Someone is interested...

...with lower label in lexicographic order

Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void unlock() {  
        flag[i] = false; ➔ No longer interested  
    }  
}
```

No Deadlock

- There is always one thread with earliest label
- Ties are impossible (why?)

First-Come-First-Served

- If $D_A \rightarrow D_B$ then A's label is earlier
 - $\text{write}_A(\text{label}[A]) \rightarrow \text{read}_B(\text{label}[A]) \rightarrow \text{write}_B(\text{label}[B]) \rightarrow \text{read}_B(\text{flag}[A])$
- So B is locked out while $\text{flag}[A] == \text{true}$

```
class Bakery implements Lock {
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while ((∃k | flag[k]) && (label[i], i) > (label[k], k)) {}
    }
}
```

Mutual Exclusion

- Suppose A and B in CS together
- Suppose A has earlier label
- When B entered, it must have seen
flag[A] == false or **label[A] > label[B]**

```
class Bakery implements Lock {  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0],...,label[n-1])+1;  
        while (( $\exists$ k | flag[k]) && (label[i],i) > (label[k],k)) {}  
    }  
}
```

Mutual Exclusion

- Labels are strictly increasing so B must have seen $\text{flag}[A] == \text{false}$
 - $\text{labeling}_B \rightarrow \text{read}_B(\text{flag}[A]) \rightarrow \text{write}_A(\text{flag}[A]) \rightarrow \text{labeling}_A$
- Which contradicts the assumption that A has an earlier label

Deep Philosophical Question

- The Bakery Algorithm is
 - Succinct
 - Elegant
 - Fair
- Q: So why isn't it practical?
- A: Well, you have to read n distinct variables

Shared Memory

- Shared read/write memory locations called **registers** (historical reasons)
- Come in different flavors
 - Multi-Reader-Single-Writer (**flag[]**)
 - Multi-Reader-Multi-Writer (**victim[]**)
 - Not interesting: SRMW and SRSW

Theorem

- At least n MRSW (multi-reader/single-writer, like `flag[]`) registers are needed to solve deadlock-free mutual exclusion

Summary of Lecture

- In the 1960's many **incorrect** solutions to lockout-free mutual exclusion using RW-registers were published...
- Today we know how to solve FIFO n thread mutual exclusion using $2n$ RW-registers